

# **EEL 6764 Principles of Computer Architecture**

## **Instruction-Level Parallelism**

Dr Hao Zheng  
Dept. of Comp Sci & Eng  
U of South Florida

# Instruction Level Parallelism

# Instruction-Level Parallelism

- Overlap executions of multiple instructions
- When exploiting instruction-level parallelism, goal is to maximize CPI
  - Pipeline CPI =
$$\frac{\text{Ideal pipeline CPI}}{\text{+ Structural stalls} + \text{+ Data hazard stalls} + \text{+ Control stalls}}$$
- Two approaches to exploiting ILP
  - Static scheduling
  - Dynamic scheduling

# Dependences

- Parallelism with **basic block** is limited
  - Typical size of basic block = 3-6 instructions
  - Must optimize across branches
- Increase ILP – Loop-Level Parallelism
  - Unroll loop statically or dynamically
  - Use SIMD (vector processors and GPUs)
- Challenges: Data dependency
  - Instruction  $j$  is data dependent on instruction  $i$  if
    - Instruction  $i$  produces a result that may be used by instruction  $j$
    - Instruction  $j$  is data dependent on instruction  $k$  and instruction  $k$  is data dependent on instruction  $i$
- Dependent instructions cannot be executed simultaneously

# Data Dependence

- Dependencies are a property of programs
- Pipeline organization determines if dependence is detected and if it causes a stall
- **Data dependence** conveys:
  - Possibility of a hazard
  - Order in which results must be calculated
  - Upper bound on exploitable instruction level parallelism
- Dependencies that flow through memory locations are difficult to detect
  - SD R1, 45(R5)
  - LD R2, 90(R7)

# Name Dependence

- Two instructions use the same name but no flow of information
  - Not a true data dependence, *but is a problem when reordering instructions*
  - **Antidependence**: instruction j writes a register or memory location that instruction i reads
    - Initial ordering (i before j) must be preserved
  - **Output dependence**: instruction i and instruction j write the same register or memory location
    - Ordering must be preserved
- To resolve, use **renaming** techniques

# Data Hazards

- A **hazard** exists between two dependent instructions that are close in a program such that overlapping their executions would change the order of accesses to the operand involved in the dependence.
- **Data Hazards**
  - Read after write (RAW)
  - Write after write (WAW)
  - Write after read (WAR)

# Control Dependences

- **Control Dependence** determines ordering of instruction  $i$  with respect to a branch instruction
  - Instruction control-dependent on a branch cannot be moved before the branch so that its execution is no longer controlled by the branch
  - An instruction not control-dependent on a branch cannot be moved after the branch so that its execution is controlled by the branch

```
if p1 { s1; } // s1 control-dependent on p1  
if p2 { s2; } // s2 control-dependent on p2
```



# Control Dependence – Examples

## Example 1:

ADD x2, x3, x4

BEQ x2, x0, L

LD x1, 0(x2)

L: ...

→ LD cannot be move above branch

→ OR instruction dependent on ADD and SUB

→ Cannot move SUB before the branch

## Example 2:

ADD x1, x2, x3

BEQ x4, x0, L

SUB x1, x5, x6

L: ...

OR x7, x1, x8

→ Reorder instruction if it does not change program semantics

# Control Dependence – Examples (1)

## Example 3

ADD x1, x2, x3

BEQ x12, x0, skip

SUB x4, x5, x6

ADD x5, x4, x9

skip: ...

OR x7, x8, x9

- Reorder instruction if it does not change program semantics
- Example 3: ok to move sub above BEQ
  - if x4 is not used after label “skip”

## **3.2 Compiler Techniques for Exposing ILP**

# Pipeline Scheduling

- Separate dependent instruction from the source instruction by the **pipeline latency** of the source instruction
  - Pipeline latency – cycles to separate two dependent instructions to avoid hazards
- By inserting NOP (bubbles) – pipeline stalls
- By inserting independent instructions - **Scheduling**
  - Statically
  - Dynamically
- Scheduling relies on
  - ILP in programs
  - latencies of functional units

# Static Scheduling

→ Transform and re-arrange the code to reduce pipeline stalls

→ Example:

```
for (i=999; i>=0; i=i-1)  
    x[i] = x[i] + s;
```

Assume branch delay = 1 cycle

Instruction producing result	Instruction using result	Latency in clock cycles
FP ALU op	Another FP ALU op	3
FP ALU op	Store double	2
Load double	FP ALU op	1
Load double	Store double	0

# Static Scheduling (1)

```
Loop: fld    f0, 0(x1)
      stall
      fdd.d  f4, f0, f2
      stall
      stall
      fsd    f4, 0(x1)
      addi   x1, x1, -8
      bne    x1, x2, Loop
```

Instruction producing result	Instruction using result	Latency in clock cycles
FP ALU op	Another FP ALU op	3
FP ALU op	Store double	2
Load double	FP ALU op	1
Load double	Store double	0

# Static Scheduling (2)

```
Loop: fld    f0, 0(x1)
      addi   x1, x1, -8
      fdd.d  f4, f0, f2
      stall
      stall
      fsd    f4, 8(x1)
      bne    x1, x2, Loop
```

- ILP in basic blocks is limited
- Loop unrolling to increase ILP

Instruction producing result	Instruction using result	Latency in clock cycles
FP ALU op	Another FP ALU op	3
FP ALU op	Store double	2
Load double	FP ALU op	1
Load double	Store double	0

# Loop Unrolling

- Unroll by a factor of 4 (assume # elements is divisible by 4)
- Eliminate unnecessary instructions

```
Loop:   fld      f0, 0(x1)
        fadd.d   f4, f0, f2
        fsd      f4, 0(R1) ;drop DADDUI & BNE
        fld      f6, -8(x1)
        fadd.d   f8, f6, f2
        fsd      f8, -8(x1) ;drop addi & BNE
        fld      f0, -16(x1)
        fadd.d   f12, f0, f2
        fsd      f12, -16(x1) ;drop addi & BNE
        fld      f14, -24(x1)
        fadd.d   f16, f14, f2
        fsd      f16, -24(x1)
        addi     x1, x1, -32
        bne     x1, x2, Loop
```

Run in 26 cycles  
without scheduling

Number of live registers  
increase vs. original loop

- Avoid dependence
- Increase ILP



# Loop Unrolling/Pipeline Scheduling

→ Pipeline schedule the unrolled loop

```
Loop:  fld      f0, 0(x1)
       fld      f6, -8(x1)
       fld      f0, -16(x1)
       fld      f14, -24(x1)
       fadd.d    f4, f0, f2
       fadd.d    f8, f6, f2
       fadd.d    f12, f0, f2
       fadd.d    f16, f14, f2
       fsd      f4, 0(x1)
       fsd      f8, -8(x1)
       fsd      f12, 16(x1)
       fsd      f16, 8(x1)
       addi     x1, x1, -32
       bne     x1, x2, Loop
```

⇒ Run in 14 cycles after scheduling

⇒ 3.5 cycles per iteration

# Dealing with Unknown Loop Bound

- Number of iterations =  $n$ , *//  $n$  is unknown*
- Goal: make  $k$  copies of the loop body
- Solution: generate pair of loops:
  - First executes  $n \bmod k$  times of original loop body
  - Second executes  $n / k$  times of unrolled loop body by  $k$  times.
  - “Strip mining”

## Limitations of Loop unrolling

- Cross loop iteration dependence
- Code size increase
- Register shortfall

## **3.4 Dynamic Scheduling**

# Dynamic Scheduling (Section 3.4, 3.5)

- In-Order pipeline
  - Issue, execution, and completion of instructions follow program order
- Pipeline stalls if there is data dependence
- But, independent instructions are stalled too.

<b>DIV.D</b>	<b>F0, F2, F4</b>
<b>ADD.D</b>	<b>F10 F0, F8</b>
<b>SUB.D</b>	<b>F12, F8, F14</b>

- Static scheduling may not identify some dependences
  - e.g. dependencies through memory addresses

# Dynamic Scheduling (Section 3.4, 3.5)

- Hardware rearrange order of instructions
  - To reduce stalls while maintaining data flow and exception
- Advantages:
  - Compiler doesn't need to have knowledge of microarchitecture
  - Handles cases where dependencies are unknown at compile time
  - Tolerate unpredictable delays in memory accesses, FP operations, etc.
- Disadvantage:
  - Substantial increase in hardware complexity
  - Complicates exceptions

# Dynamic Scheduling

- ID stage is divided into
  - **Issue** – decode instruction, and check for S-hazards
  - **Read Operands** – wait until no D-hazards, then read operands
- Instructions pass **Issue** in-order
- They are stalled or bypass each other in **Read Operands**
  - Enter **Execute** when operands are ready
  - Enter **Execute** out-of-order

# Dynamic Scheduling

- Dynamic scheduling implies:
  - In-order issue
  - Out-of-order execution
  - Out-of-order completion
- Creates the possibility for WAR and WAW hazards

<b>DIV.D</b>	<b>F0, F2, F4</b>
<b>ADD.D</b>	<b>F6, F0, F8</b>
<b>SUB.D</b>	<b>F8, F10, F14</b>
<b>MUL.D</b>	<b>F6, F10, F8</b>

- WAR/WAW hazards can be avoided by **register renaming**

# Dynamic Scheduling

- Tomasulo's Approach
  - Tracks when operands are available
    - To eliminate **RAW** hazards
  - Introduces **register renaming** in hardware
    - To eliminate **WAW** and **WAR** hazards



# Register Renaming

antidependence via **F8**



DIV.D	F0, F2, F4
ADD.D	<b>F6</b> , F0, <b>F8</b>
sd	f6, 0(x1)
SUB.D	<b>F8</b> , F10, F14
MUL.D	<b>F6</b> , F10, F8

Output dependence via **F6**

# Register Renaming

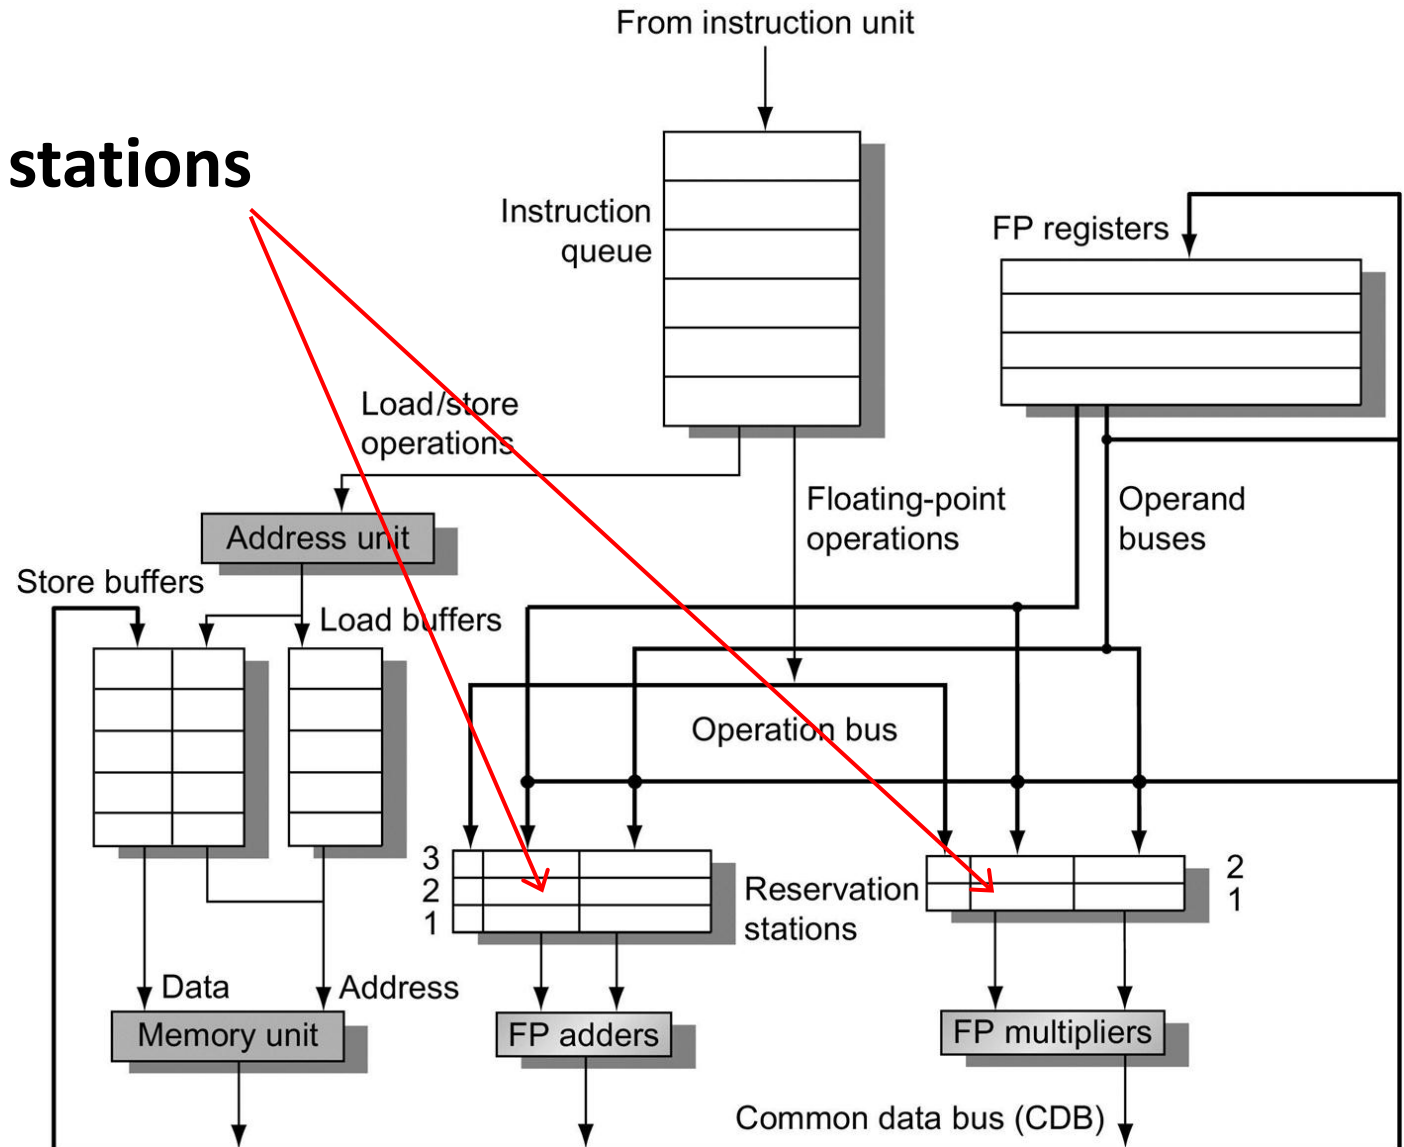
→ Example: assume two additional registers **S** and **T**

<b>DIV.D</b>	<b>F0, F2, F4</b>	
<b>ADD.D</b>	<b>S, F0, F8</b>	<b>; F6 -&gt; S</b>
<b>sd</b>	<b>S, 0(x1)</b>	<b>; F6 -&gt; S</b>
<b>SUB.D</b>	<b>T, F10, F14</b>	<b>; F8 -&gt; T</b>
<b>MUL.D</b>	<b>F6, F10, T</b>	<b>; F8 -&gt; T</b>

→ Now only **RAW** hazards remain, which must be strictly ordered

# Tomasulo's Algorithm – Structure

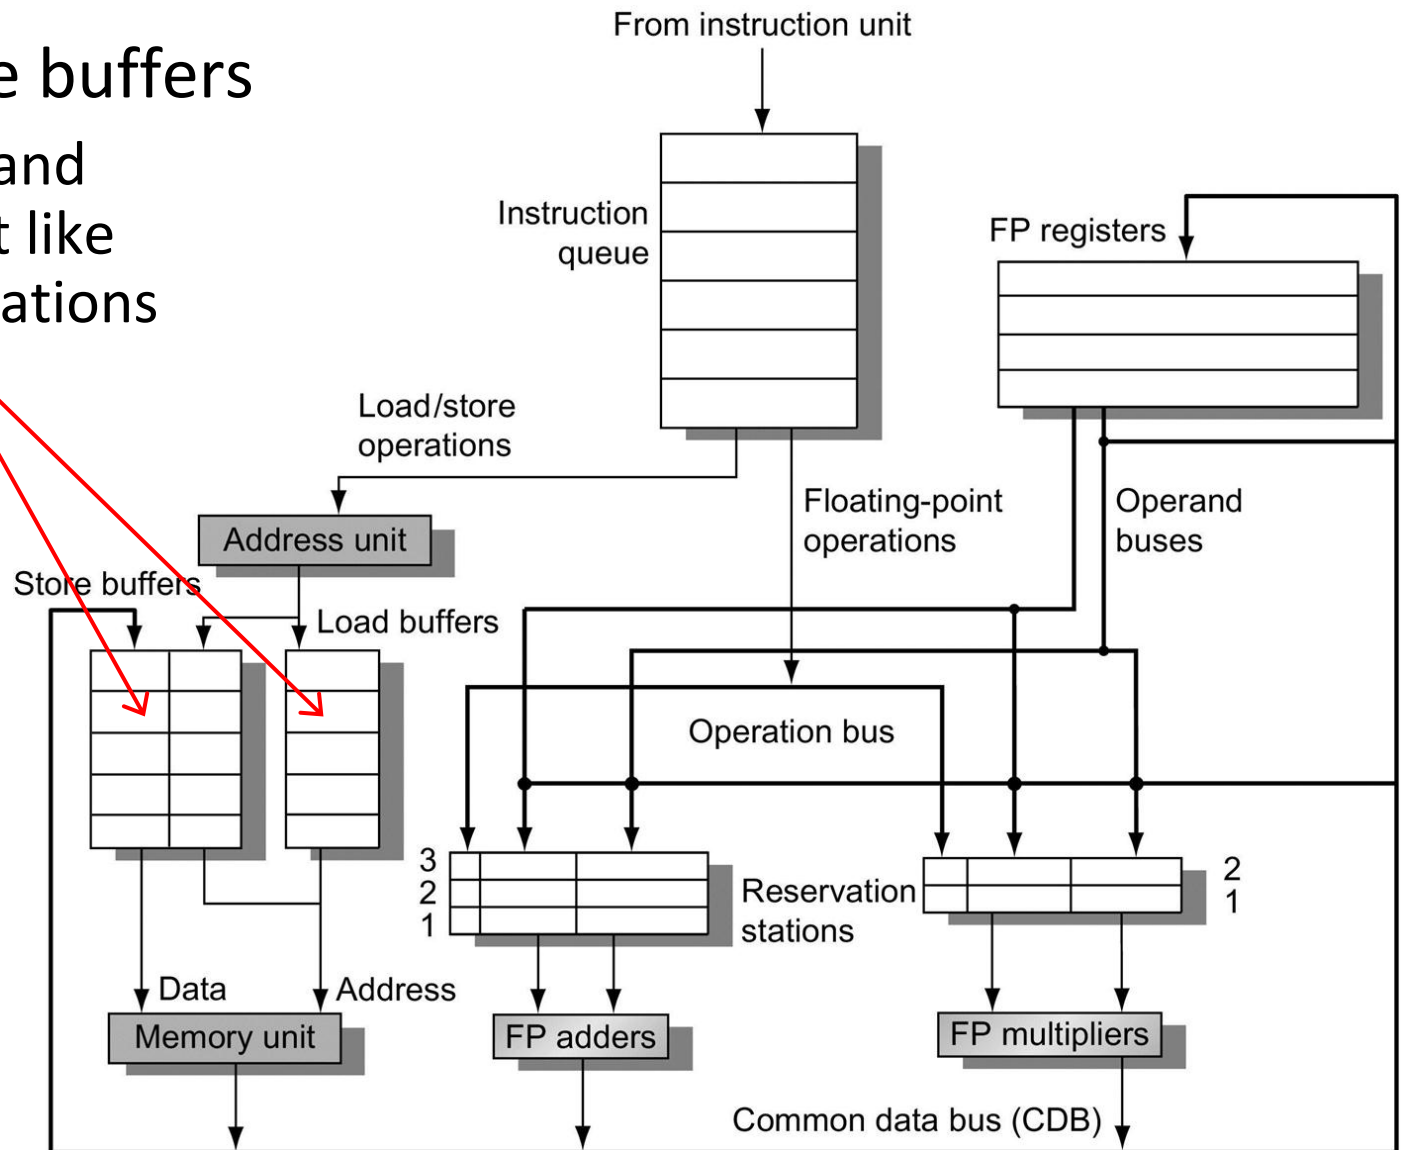
## Reservation stations



# Tomasulo's Algorithm – Structure

## Load and store buffers

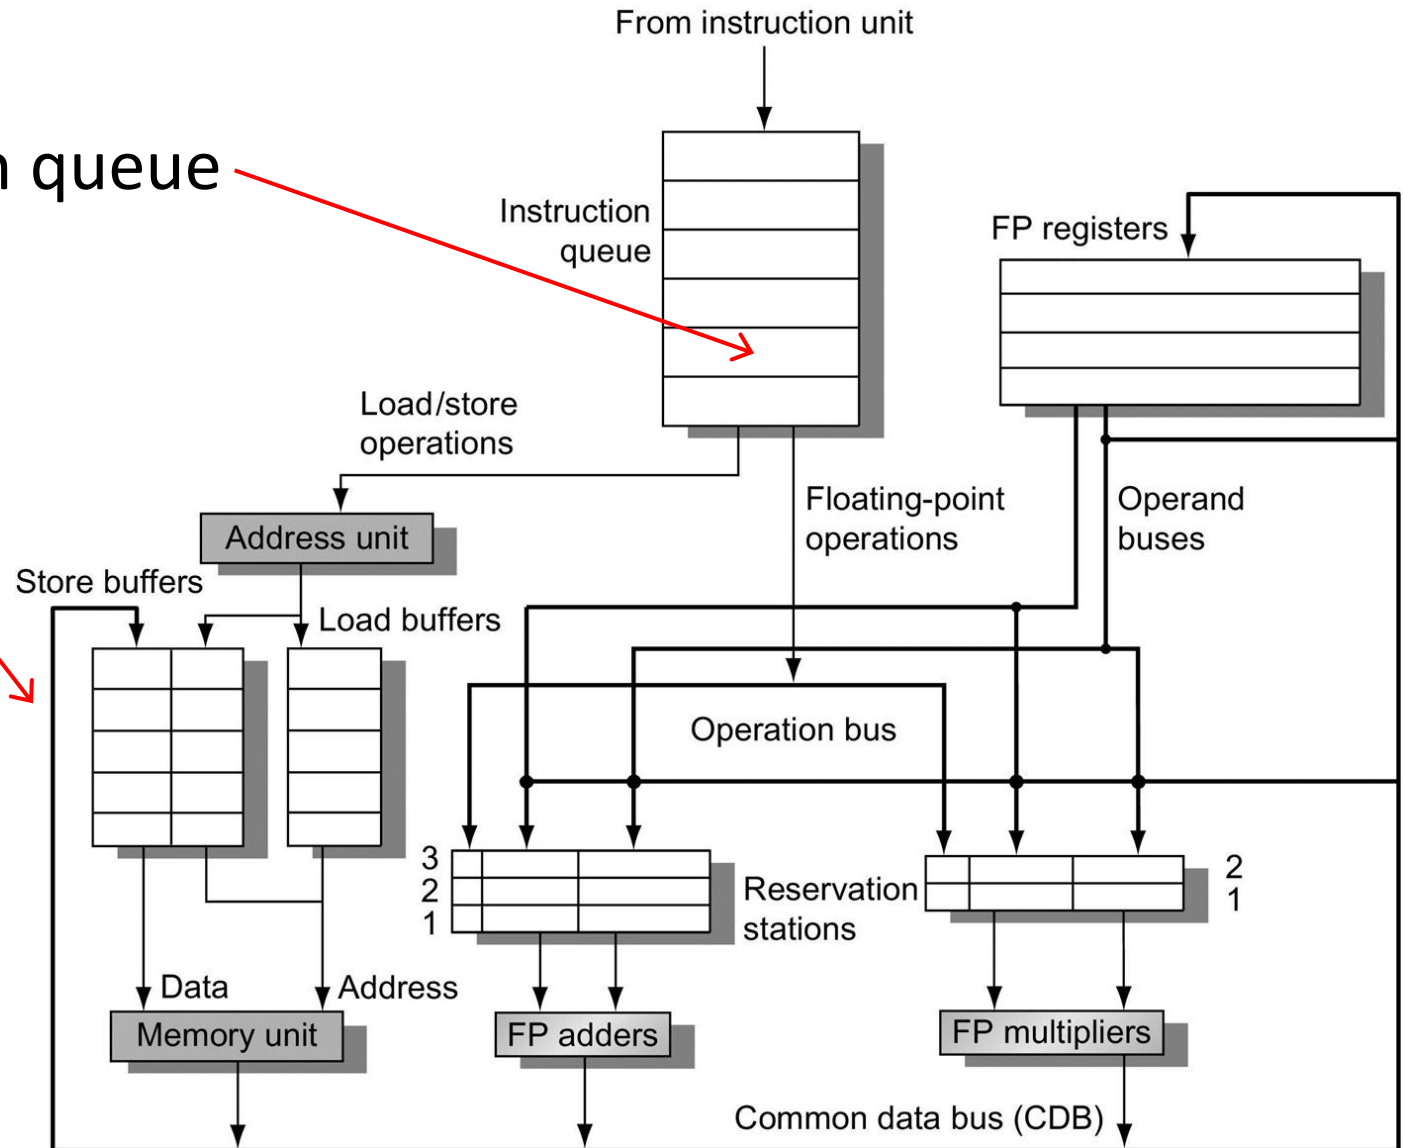
→ Contain data and addresses, act like reservation stations



# Tomasulo's Algorithm – Structure

→ Instruction queue

→ CDB

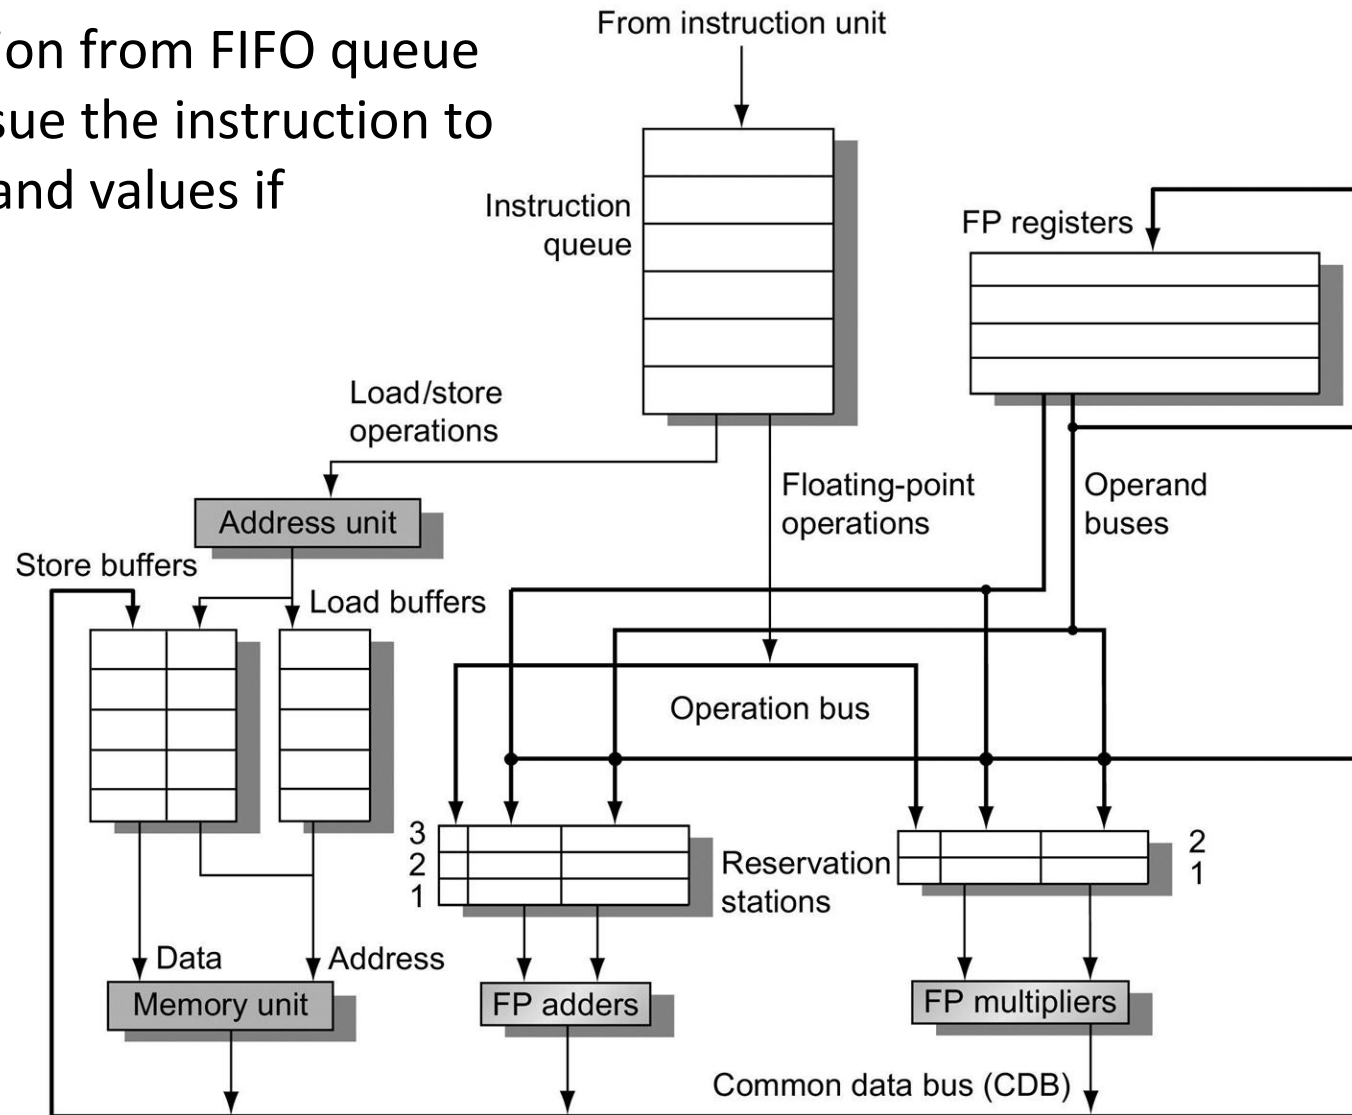


# Register Renaming By Reservation Stations

- Reservation stations (RS) buffer
  - The instruction opcode
  - Buffered operand values (when available)
  - Or, reservation station number of instruction providing the operand values (**producer instruction**)
- RS fetches and buffers an operand as soon as it becomes available (not necessarily involving register file)
- Pending instructions designate the RS to which they will send their output
- Result values broadcast on a result bus, common data bus (CDB)
- Only the last output updates the register file
- As instructions are issued, the register specifiers are renamed with the reservation station
- May be more reservation stations than registers

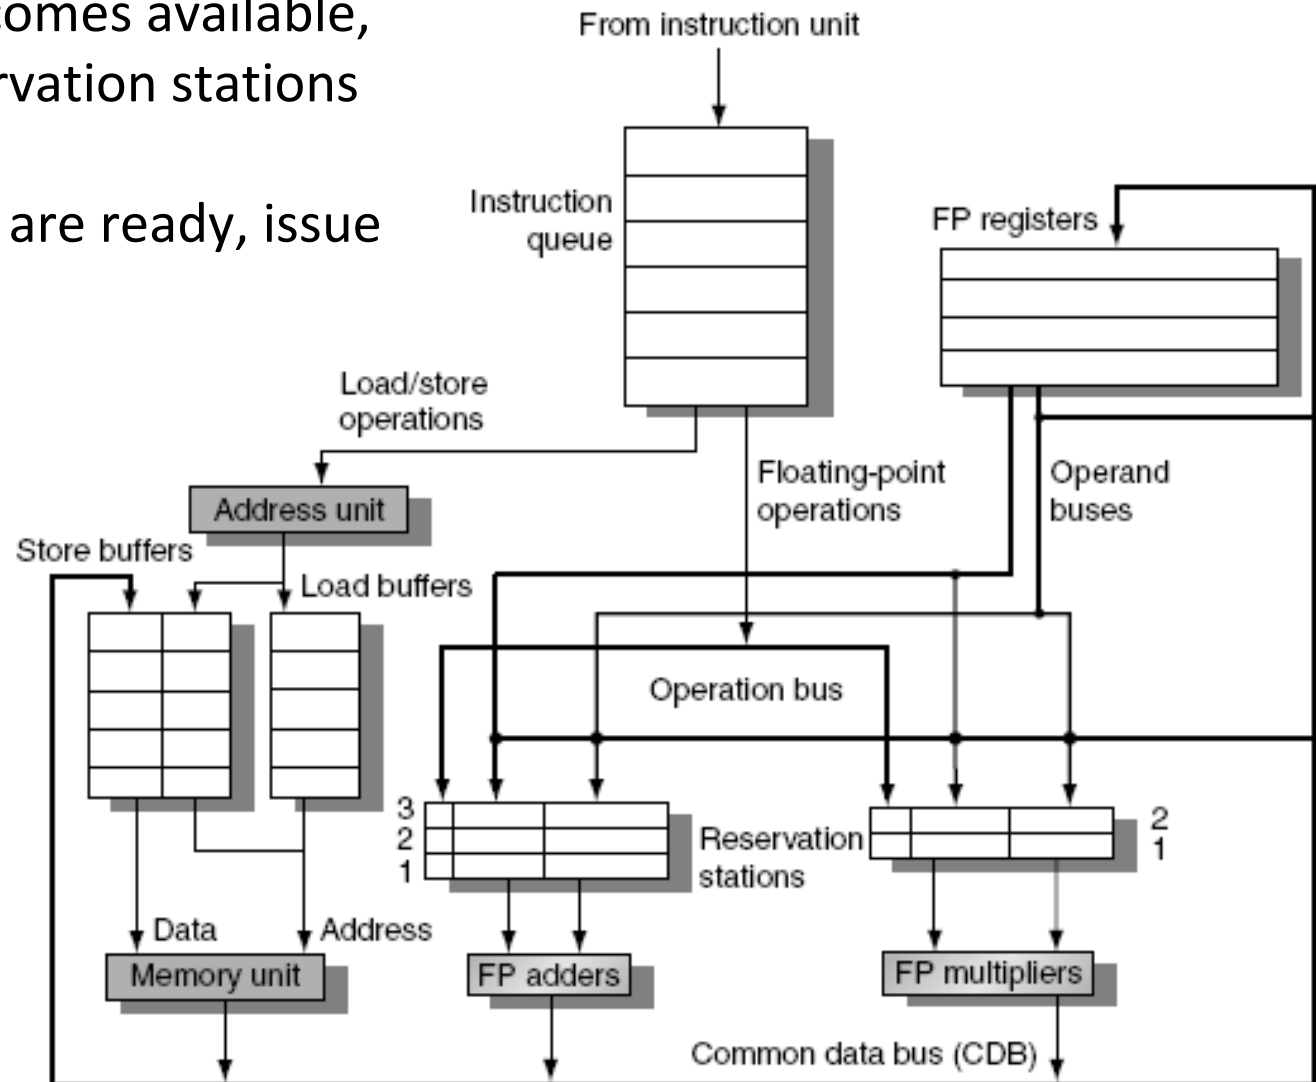
# Tomasulo's Algorithm – Issue

- Get next instruction from FIFO queue
- If available RS, issue the instruction to the RS with operand values if available
- If operand values not available, stall the instruction
- Renaming in this step



# Tomasulo's Algorithm – Execute

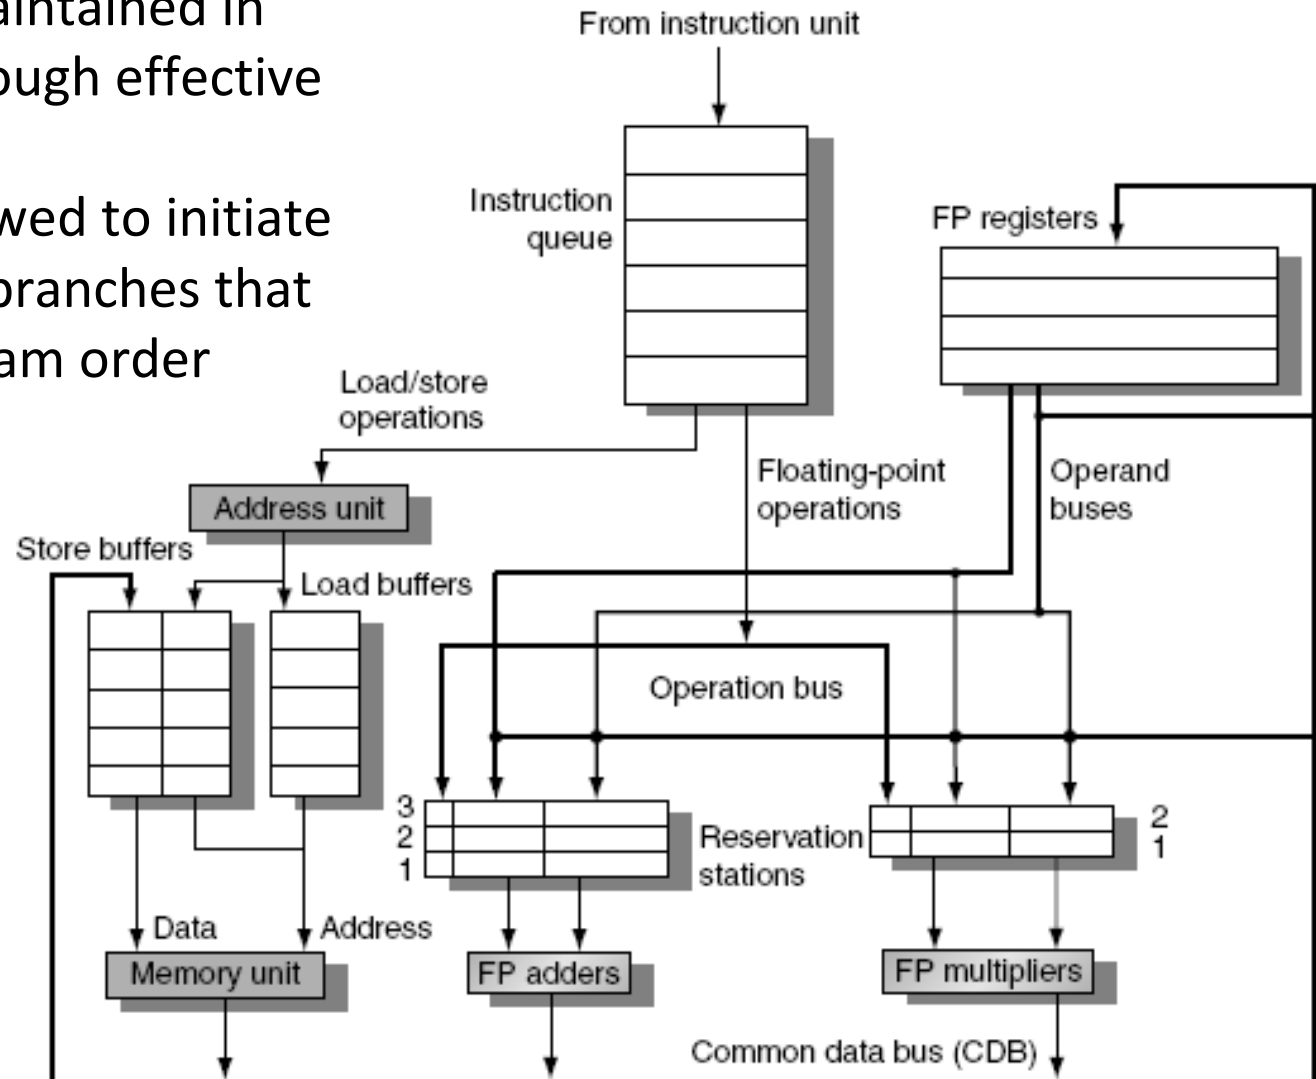
- When operand becomes available, store it in any reservation stations waiting for it
- When all operands are ready, issue the instruction





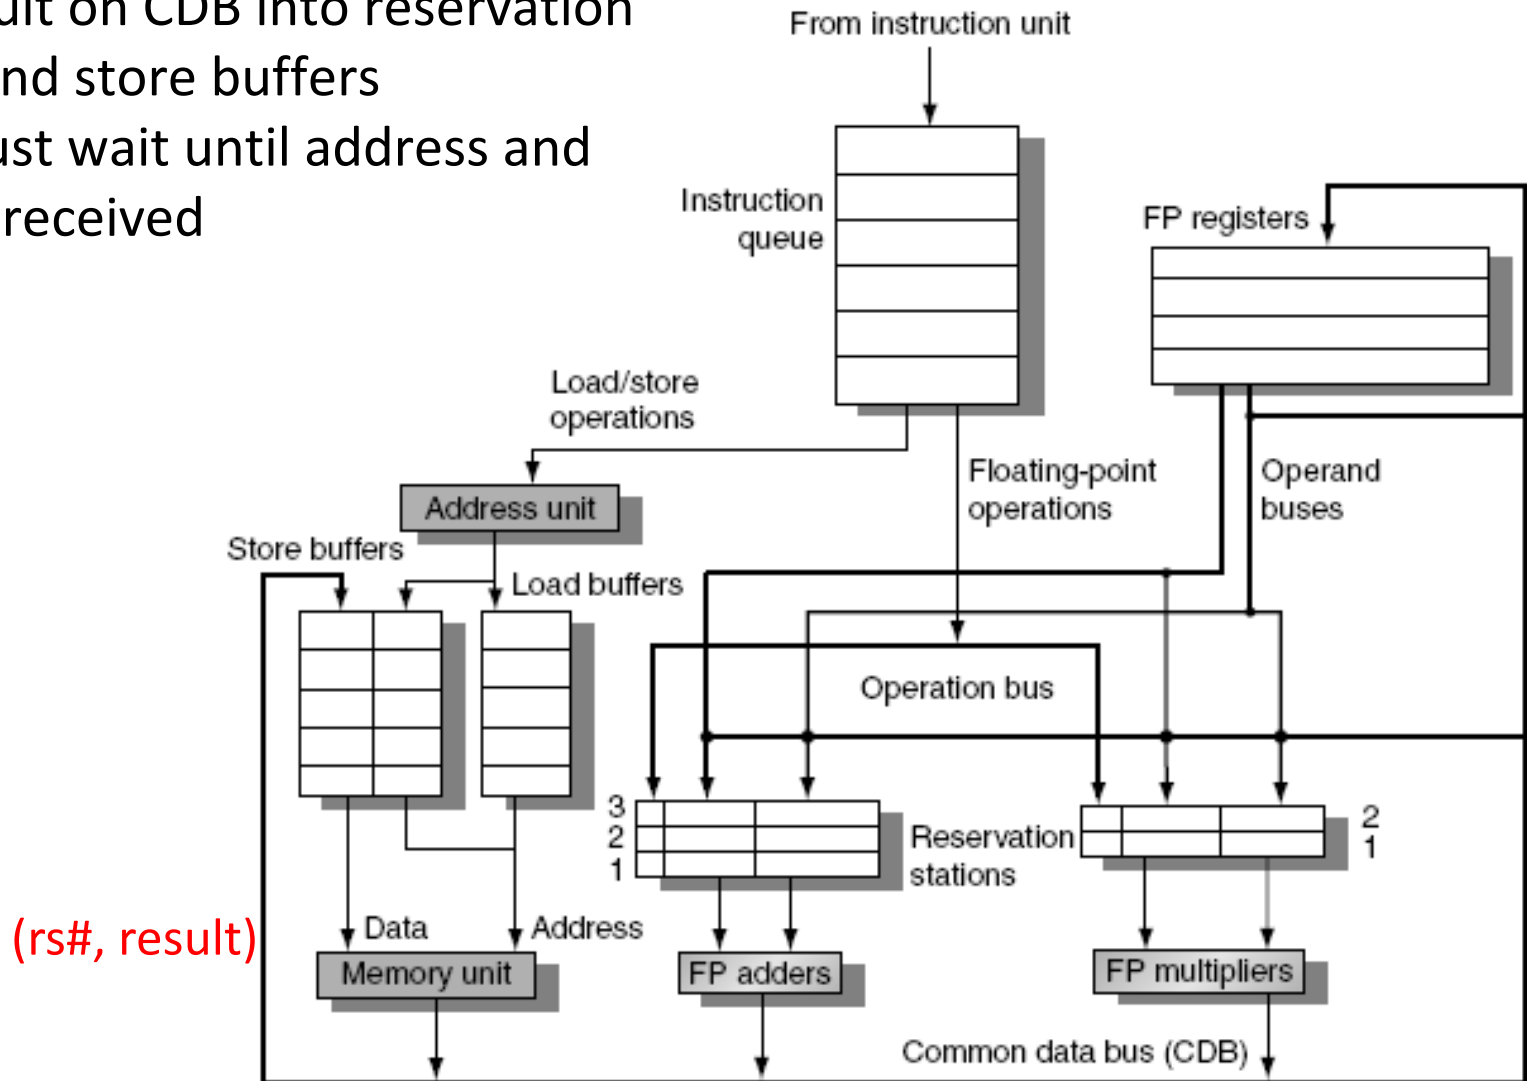
# Tomasulo's Algorithm – Execute (1)

- Loads and store maintained in program order through effective address
- No instruction allowed to initiate execution until all branches that proceed it in program order have completed



# Tomasulo's Algorithm – Write Result

- Write result on CDB into reservation stations and store buffers
- Stores must wait until address and value are received



# Tomasulo's Algorithm – Some Notes

- Each RS has an unique ID
- Each register is tagged with RS#
  - Will be updated by instruction in RS#
- Each RS, store buffer tagged with RS# for unavailable operands
  - Tags refer to FUs or buffers that produce operands
  - Tags = producer ID
- FUs or memory (producer) put (RS#, results) on CDB
- RS, registers, store buffers (consumer) monitor the CDB,
  - Get the data where its tag matches tag on CDB
- Forwarding through broadcast on CDB

# Example

Instruction		Instruction status		
		Issue	Execute	Write result
fld	f6,32(x2)	✓	✓	✓
fld	f2,44(x3)	✓	✓	
fmul.d	f0,f2,f4	✓		
fsub.d	f8,f2,f6	✓		
fdiv.d	f0,f0,f6	✓		
fadd.d	f6,f8,f2	✓		

Reservation stations							
Name	Busy	Op	Vj	Vk	Qj	Qk	A
Load1	No						
Load2	Yes	Load					44 + Regs[x3]
Add1	Yes	SUB		Mem[32 + Regs[x2]]	Load2		
Add2	Yes	ADD			Add1	Load2	
Add3	No						
Mult1	Yes	MUL		Regs[f4]	Load2		
Mult2	Yes	DIV		Mem[32 + Regs[x2]]	Mult1		

Register status									
Field	f0	f2	f4	f6	f8	f10	f12	...	f30
Qi	Mult1	Load2		Add2	Add1	Mult2			

# Reservation Station Fields

## Instruction status

- **Op** – inst opcode
- **Qj, Qk** – RS that produce the operands
- **Vj, Vk** – operand values
- **A** – hold memory address for load/store
- **Qi** – RS whose output should be stored into this register

## Reservation stations

Name	Busy	Op	Vj	Vk	Qj	Qk	A
Load1	No						
Load2	Yes	Load					44 + Regs[x3]
Add1	Yes	SUB		Mem[32 + Regs[x2]]	Load2		
Add2	Yes	ADD			Add1	Load2	
Add3	No						
Mult1	Yes	MUL		Regs[f4]	Load2		
Mult2	Yes	DIV		Mem[32 + Regs[x2]]	Mult1		

## Register status

Field	f0	f2	f4	f6	f8	f10	f12	...	f30
Qi	Mult1	Load2		Add2	Add1	Mult2			

## Instruction status

### Instruction

f1d	f6,32(x2)
f1d	f2,44(x3)
fmul.d	f0,f2,f4
fsub.d	f8,f2,f6
fdiv.d	f0,f0,f6
fadd.d	f6,f8,f2

### Write result

## Latencies of FUs

- Load = 1 cycle
- Add/sub = 2 cycles
- Mult = 6 cycles
- Div = 12 cycles

## Reservation stations

Name	Busy	Op	Vj	Vk	Qj	Qk	A
Load1							
Load2							
Add1							
Add2							
Add3							
Mult1							
Mult2							

## Register status

Field	f0	f2	f4	f6	f8	f10	f12	...	f30
Qi									

Instruction	Instruction status		
	Issue	Execute	Write result
f1d f6,32(x2)			
f1d f2,44(x3)			
fmul.d f0,f2,f4			
fsub.d f8,f2,f6			
fdiv.d f0,f0,f6			
fadd.d f6,f8,f2			

Reservation stations							
Name	Busy	Op	Vj	Vk	Qj	Qk	A
Load1							
Load2							
Add1							
Add2							
Add3							
Mult1							
Mult2							

Register status									
Field	f0	f2	f4	f6	f8	f10	f12	...	f30
Qi									

# A Loop Example

```
Loop: fld      f0, 0(x1)
      fmul     f4, f0, f2
      fsd      f4, 0(x1)
      addi     x1, x1, -8
      bne      x1, x2, Loop
```

f#: FP registers

x#: integer registers

Assume predict-taken for loop unrolling



Instruction		Instruction status			
		From iteration	Issue	Execute	Write result
fld	f0,0(x1)	1	✓	✓	
fmul.d	f4,f0,f2	1	✓		
fsd	f4,0(x1)	1	✓		
fld	f0,0(x1)	2	✓	✓	
fmul.d	f4,f0,f2	2	✓		
fsd	f4,0(x1)	2	✓		

Reservation stations							
Name	Busy	Op	Vj	Vk	Qj	Qk	A
Load1	Yes	Load					Regs[x1] + 0
Load2	Yes	Load					Regs[x1] − 8
Add1	No						
Add2	No						
Add3	No						
Mult1	Yes	MUL		Regs[f2]	Load1		
Mult2	Yes	MUL		Regs[f2]	Load2		
Store1	Yes	Store	Regs[x1]			Mult1	
Store2	Yes	Store	Regs[x1] − 8			Mult2	

Register status									
Field	f0	f2	f4	f6	f8	f10	f12	...	f30
Qi	Load2		Mult2						

# Ordering of Load and Store

- If load and store accesses different memory, they can be re-ordered
- Otherwise, reordering causes WAW/RAW hazards

sd	f1, XXX
ld	f2, XXX

ld	f1, XXX
sd	f2, XXX

sd	f1, XXX
sd	f2, XXX

- Consider **load** is next to issue
  - Computer effective address for **load**
  - Check the **A** field in all **store buffers** for match
  - Hold **load** if a match is found
- Store is handled similarly
  - Check both **load** and **store** buffers

# Tomasulo's Algorithm – Summary

- Rename registers to RS
  - To eliminates WAW and WAR hazards
- Buffer RS # for operands currently unavailable
  - To eliminate RAW hazards
- Tolerate unpredictable delays of mem hierarchy
  - Allow instructions to execute while waiting for cache misses
- Achieve high performance w/o compiler optimization
- HW complexity increases substantially
  - Increase area and power

## **3.6 Hardware-Based Speculation**

# Hardware-Based Speculation

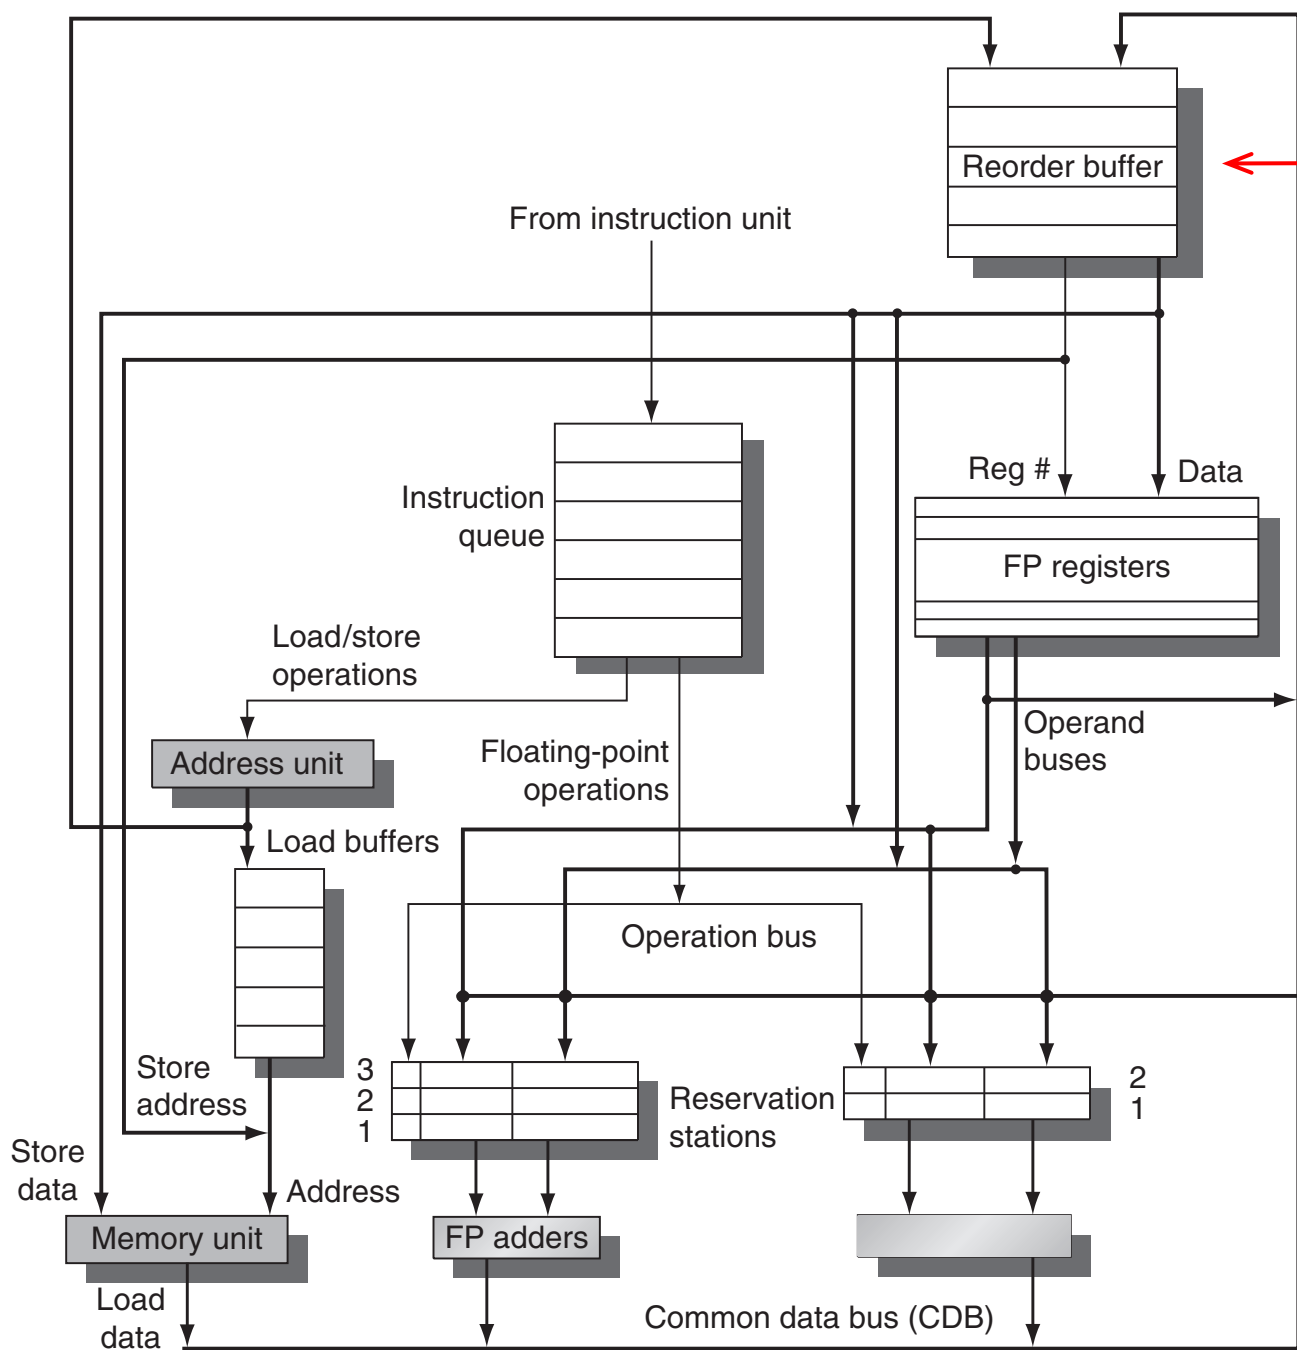
- Combine **dynamic scheduling** and **branch prediction**
- **Execute** instructions along predicted execution paths but only **commit** the results if prediction was correct
  - Speculative execution
- **Instruction commit**: allowing an instruction to update the register/mem when instruction is no longer speculative
  - i.e. branch prediction is correct.
- Instruction **execute out-of-order**, **commit in-order**
- Need an additional piece of hardware to prevent any irrevocable action until an instruction commits
  - I.e. updating state or taking an execution

# Reorder Buffer

- Reorder buffer (ROB) – holds the result of instruction between completion and commit
- Four fields:
  - Instruction type: branch/store/register
  - Destination field: register number/memory address
  - Value field: output value
  - Ready field: completed execution?
- Source operands are in
  - Registers  $\leftarrow$  committed instructions
  - ROB  $\leftarrow$  instructions that complete execution before commit
- Modify reservation stations:
  - Operand sources ( $Q_j$  or  $Q_k$ ) are now ROB entries instead of  $RS\#$


# Reorder Buffer

- ROB and store buffer are merged
- Values to registers/memory are not written until an instruction commits
- On mis-prediction:
  - Speculated entries in ROB are cleared
- Exceptions:
  - Not recognized until it is ready to commit




Combined w  
store buffers –  
A queue





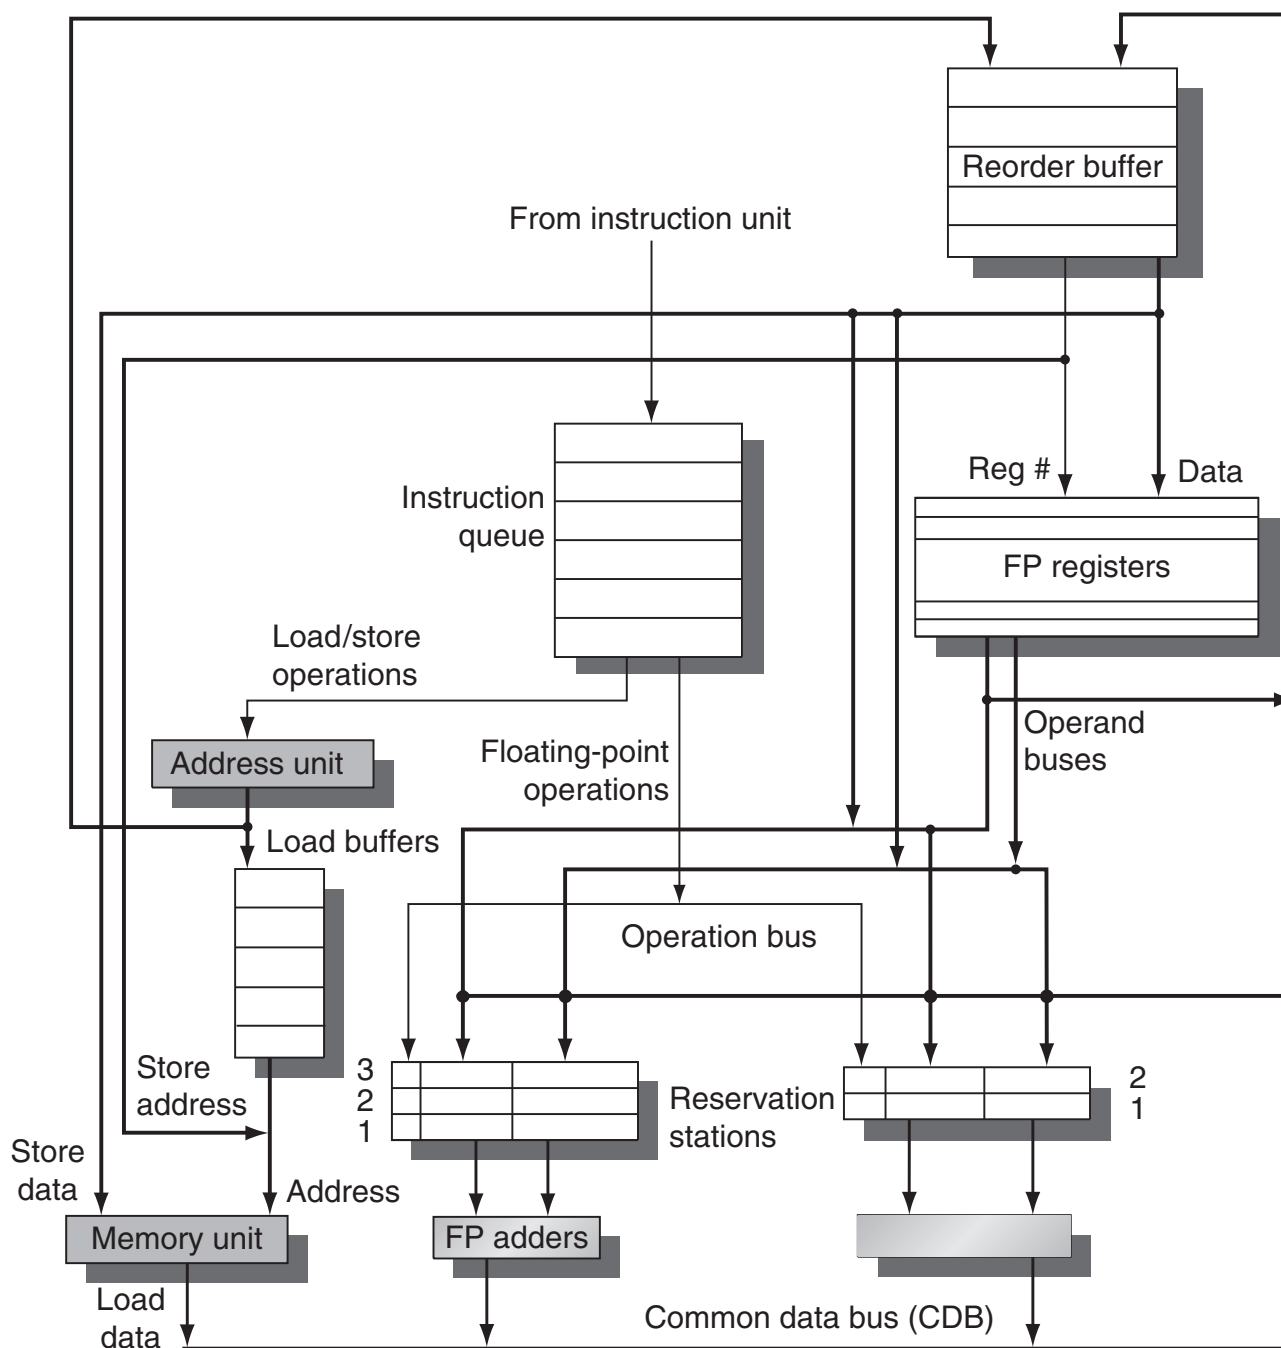
- **Issue** an instruction if there is an empty RS and ROB. Stall otherwise.
- **Send the operands to RS if they are available or ROB# if not.**
- **Tag the instruction result with the ROB# for the destination register**

(ROB#, value)




- **Issue**
- **Execute** – start after both operands are available.
- **Results are tagged with ROB#**

(ROB#, value)



- **Issue**
- **Execute** – start after both operands are available.
- **Write Result** – result put on CDB and goes to ROB and RS
- Any RS or ROB entry expecting a result with ROB# will grab value when <ROB#, value> is put on CDB

(ROB#, value)



- **Issue**
- **Execute**
- **Write Result**
- **Commit** – write result of instruction at the head of ROB
- If the instruction is NOT branch, update M/Reg with the **value** in ROB
- If the head of ROB is a **mis-predicted** branch, flush ROB, and restart from the branch target.
- Otherwise, commit as normal.

(ROB#, value)

Reorder buffer						
Entry	Busy	Instruction		State	Destination	Value
1	No	fld	f6,32(x2)	Commit	f6	Mem[32 + Regs[x2]]
2	No	fld	f2,44(x3)	Commit	f2	Mem[44 + Regs[x3]]
3	Yes	fmul.d	f0,f2,f4	Write result	f0	#2 × Regs[f4]
4	Yes	fsub.d	f8,f2,f6	Write result	f8	#2 − #1
5	Yes	fdiv.d	f0,f0,f6	Execute	f0	
6	Yes	fadd.d	f6,f8,f2	Write result	f6	#4 + #2

Reservation stations								
Name	Busy	Op	Vj	Vk	Qj	Qk	Dest	A
Load1	No							
Load2	No							
Add1	No							
Add2	No							
Add3	No							
Mult1	No	fmul.d	Mem[44 + Regs[x3]]	Regs[f4]			#3	
Mult2	Yes	fdiv.d		Mem[32 + Regs[x2]]	#3		#5	

FP register status										
Field	f0	f1	f2	f3	f4	f5	f6	f7	f8	f10
Reorder #	3						6		4	5
Busy	Yes	No	No	No	No	No	Yes	...	Yes	Yes

# A Loop Example

Reorder buffer						
Entry	Busy	Instruction		State	Destination	Value
1	No	fld	f0,0(x1)	Commit	f0	Mem[0 + Regs[x1]]
2	No	fmul.d	f4,f0,f2	Commit	f4	#1 × Regs[f2]
3	Yes	fsd	f4,0(x1)	Write result	0 + Regs[x1]	#2
4	Yes	addi	x1,x1,-8	Write result	x1	Regs[x1] - 8
5	Yes	bne	x1,x2,Loop	Write result		
6	Yes	fld	f0,0(x1)	Write result	f0	Mem[#4]
7	Yes	fmul.d	f4,f0,f2	Write result	f4	#6 × Regs[f2]
8	Yes	fsd	f4,0(x1)	Write result	0 + #4	#7
9	Yes	addi	x1,x1,-8	Write result	x1	#4 - 8
10	Yes	bne	x1,x2,Loop	Write result		

FP register status									
Field	f0	f1	f2	f3	f4	F5	f6	F7	f8
Reorder #	6								
Busy	Yes	No	No	No	Yes	No	No	...	No

# Hazards Through Memory

- No WAW/WAR hazards by
  - In-order commit to update reg/memory
- *Store updates MEM in **commit**, while load reads MEM in **execute***
- Avoid RAW hazards through MEM by
  - Not allowing a **load** to read Mem if its **A** field matches the destination field of any active ROB entry for a store, and
  - Maintaining the program order of effective address computation of a **load** wrt all earlier **stores**.

```
sd    f1, XXX
ld    f2, XXX
```

# Hardware Speculation – Summary

- Out-of-order execution, in-order commit
- On branch mis-prediction, flush ROB, and processor restart from the branch target.
  - Instructions before the branch are completed as normal
- Branch prediction has higher importance.
  - Impact of mis-prediction is higher
- Instruction exception is delayed until instruction commit
  - Allow precise exception
- WAW/WAR hazards are eliminated by in-order commit
- No RAW hazards with additional control



## **3.7 -- 3.8 Multiple Issue**

# Multiple Issue and Static Scheduling

- To achieve  $CPI < 1$ , need to complete multiple instructions per clock
- Solutions:
  - Statically scheduled superscalar processors
  - VLIW (very long instruction word) processors
  - dynamically scheduled superscalar processors

# Multiple Issue

Common name	Issue structure	Hazard detection	Scheduling	Distinguishing characteristic	Examples
Superscalar (static)	Dynamic	Hardware	Static	In-order execution	Mostly in the embedded space: MIPS and ARM, including the ARM Cortex A8
Superscalar (dynamic)	Dynamic	Hardware	Dynamic	Some out-of-order execution, but no speculation	None at the present
Superscalar (speculative)	Dynamic	Hardware	Dynamic with speculation	Out-of-order execution with speculation	Intel Core i3, i5, i7; AMD Phenom; IBM Power 7
VLIW/LIW	Static	Primarily software	Static	All hazards determined and indicated by compiler (often implicitly)	Most examples are in signal processing, such as the TI C6x
EPIC	Primarily static	Primarily software	Mostly static	All hazards determined and indicated explicitly by the compiler	Itanium

# VLIW Processors

- Package multiple operations into one instruction
- Example VLIW processor:
  - One integer instruction (or branch)
  - Two independent floating-point operations
  - Two independent memory references
- Must be enough parallelism in code to fill the available slots
  - loop unrolling is a major technique

# VLIW Processors

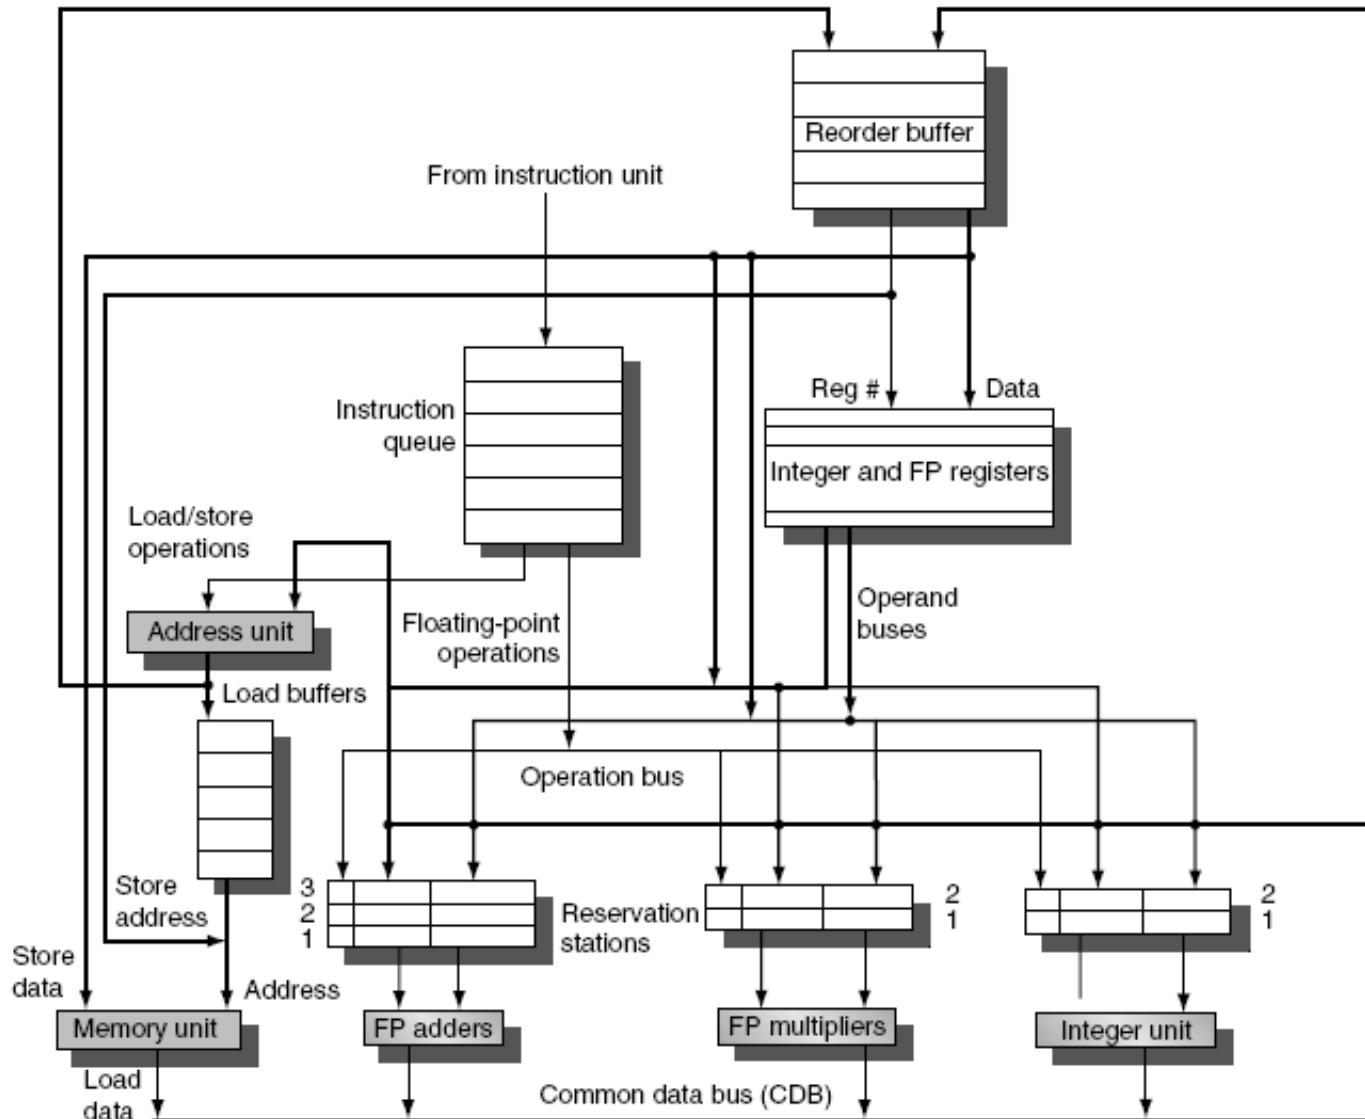
## → Disadvantages:

- Statically finding parallelism – loop unrolling
- Single code size increase
- No hazard detection hardware
- Binary code compatibility

# Dynamic Scheduling, Multiple Issue, and Speculation

- Modern microarchitectures:
  - Dynamic scheduling + multiple issue + speculation
- Two approaches:
  - Assign reservation stations and update pipeline control table in half clock cycles
    - Only supports 2 instructions/clock
  - Design logic to handle any possible dependencies between the instructions
  - Hybrid approaches
- Issue logic can become bottleneck
  - see fig. 3.22 for issue logic for (FP load + FP op)

# Overview of Design



Multiple FUs  
(integer and FP)

An instruction  
issued for each FU

Instructions issued  
in-order

RS and ROB  
updated for all  
issued instructions

# Multiple Issue

- Limit the number of instructions of a given class that can be issued in a “bundle”, and pre-allocate RS and ROB
  - I.e. on FP, one integer, one load, one store
- Examine all the dependencies among the instructions in the bundle
- If dependencies exist in bundle, encode them in reservation stations
- Also need multiple completion/commit
  - easier – dependencies among issued instructions resolved



# Example

Loop:	ld	x2, 0(x1)	;R2=array element
	addi	x2, x2, 1	;increment R2
	sd	x2, 0(x1)	;store result
	addi	x1, x1, 8	;increment pointer
	bne	x2, x3, loop	;branch if not last element

What happens when issue both?

ld	x2, 0(x1)
addi	x2, x2, 1

# Example – Mult-Issue, No Speculation

Iteration number	Instructions	Issues at clock cycle number	Executes at clock cycle number	Memory access at clock cycle number	Write CDB at clock cycle number	Comment
1	ld x2,0(x1)	1	2	3	4	First issue
1	addi x2,x2,1	1	5		6	Wait for ld
1	sd x2,0(x1)	2	3	7		Wait for addi
1	addi x1,x1,8	2	3		4	Execute directly
1	bne x2,x3,Loop	3	7			Wait for addi
2	ld x2,0(x1)	4	8	9	10	Wait for bne
2	addi x2,x2,1	4	11		12	Wait for ld
2	sd x2,0(x1)	5	9	13		Wait for addi
2	addi x1,x1,8	5	8		9	Wait for bne
2	bne x2,x3,Loop	6	13			Wait for addi
3	ld x2,0(x1)	7	14	15	16	Wait for bne
3	addi x2,x2,1	7	17		18	Wait for ld
3	sd x2,0(x1)	8	15	19		Wait for addi
3	addi x1,x1,8	8	14		15	Wait for bne
3	bne x2,x3,Loop	9	19			Wait for addi

# Example – Mult-Issue w Speculation

Iteration number	Instructions	Issues at clock number	Executes at clock number	Read access at clock number	Write CDB at clock number	Commits at clock number	Comment
1	ld x2,0(x1)	1	2	3	4	5	First issue
1	addi x2,x2,1	1	5		6	7	Wait for ld
1	sd x2,0(x1)	2	3			7	Wait for addi
1	addi x1,x1,8	2	3		4	8	Commit in order
1	bne x2,x3,Loop	3	7			8	Wait for addi
2	ld x2,0(x1)	4	5	6	7	9	No execute delay
2	addi x2,x2,1	4	8		9	10	Wait for ld
2	sd x2,0(x1)	5	6			10	Wait for addi
2	addi x1,x1,8	5	6		7	11	Commit in order
2	bne x2,x3,Loop	6	10			11	Wait for addi
3	ld x2,0(x1)	7	8	9	10	12	Earliest possible
3	addi x2,x2,1	7	11		12	13	Wait for ld
3	sd x2,0(x1)	8	9			13	Wait for addi
3	addi x1,x1,8	8	9		10	14	Executes earlier
3	bne x2,x3,Loop	9	13			14	Wait for addi

## **3.11 Multithreading**


# Limitations of ILP

- Program structure – limited ILP
- WAW/WAR hazards through memory
- Memory bandwidth and latency
  - Pipeline cannot hide latency to access off-chip cache/mem.
  - Remember “memory wall”?
- Impacts of wide issue width
  - Logic complexity ↑
  - Clock rate ↓
  - Power ↑
- Size of ROB and RS


# Multithreading

- Exploiting thread-level parallelism to improve uniprocessor throughput
- Multiprocessing and multithreading in Chapter 5
- Thread-level parallelism (TLP) exists abundantly
  - A property of applications,
  - eg. online transaction processing, SC, web applications,
- Multithreading allows multiple threads to share a processor without process switch
  - Each thread is independently controlled by OS
  - With duplicate private state (reg, PC, etc) for each thread
  - Share FUs and memory, etc

# Multithreading Hardware



# Effectiveness of SMT





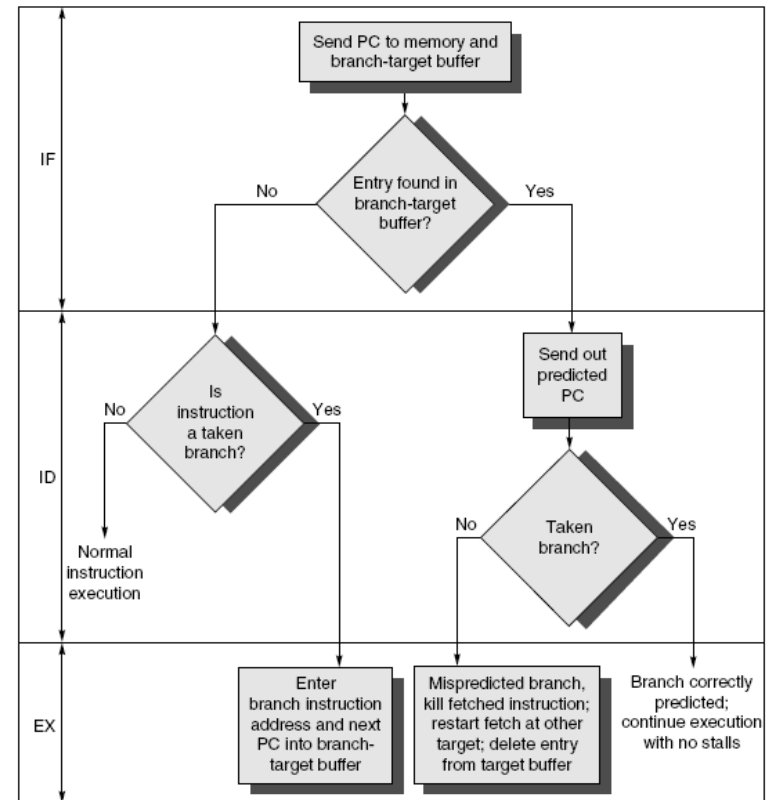
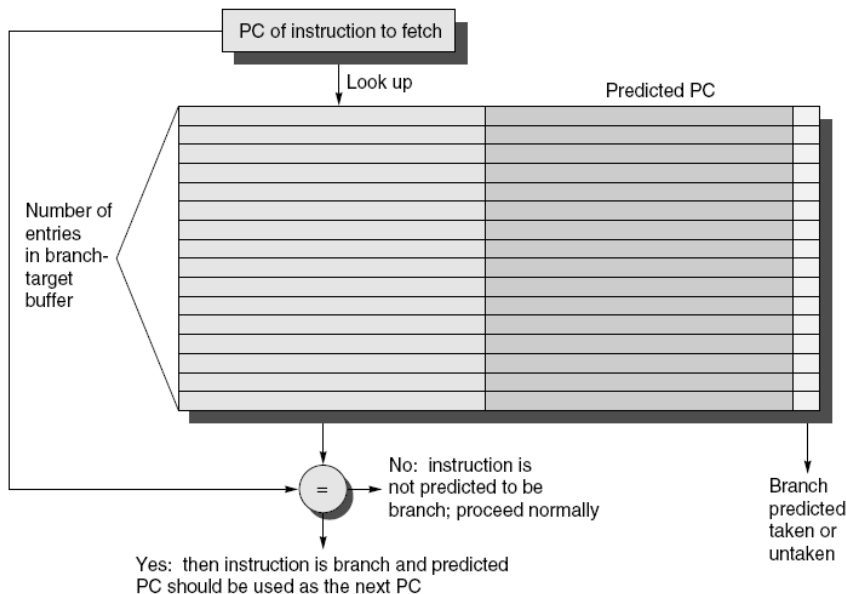
# Backup

# Branch-Target Buffer

→ Need high instruction bandwidth!

→ Branch-Target buffers

→ Next PC prediction buffer, indexed by current PC



# Branch Folding

## → Optimization:

- Larger branch-target buffer
- Add target instruction into buffer to deal with longer decoding time required by larger buffer
- “Branch folding”

# Return Address Predictor

- Most unconditional branches come from function returns
- The same procedure can be called from multiple sites
  - Causes the buffer to potentially forget about the return address from previous calls
- Create return address buffer organized as a stack

# Integrated Instruction Fetch Unit

- Design monolithic unit that performs:
  - Branch prediction
  - Instruction prefetch
    - Fetch ahead
  - Instruction memory access and buffering
    - Deal with crossing cache lines

# Register Renaming

- Register renaming vs. reorder buffers
  - Instead of virtual registers from reservation stations and reorder buffer, create a single register pool
    - Contains visible registers and virtual registers
  - Use hardware-based map to rename registers during issue
  - WAW and WAR hazards are avoided
  - Speculation recovery occurs by copying during commit
  - Still need a ROB-like queue to update table in order
  - Simplifies commit:
    - Record that mapping between architectural register and physical register is no longer speculative
    - Free up physical register used to hold older value
    - In other words: SWAP physical registers on commit
  - Physical register de-allocation is more difficult

# Integrated Issue and Renaming

- Combining instruction issue with register renaming:
  - Issue logic pre-reserves enough physical registers for the bundle (fixed number?)
  - Issue logic finds dependencies within bundle, maps registers as necessary
  - Issue logic finds dependencies between current bundle and already in-flight bundles, maps registers as necessary

# How Much?

- How much to speculate
  - Mis-speculation degrades performance and power relative to no speculation
    - May cause additional misses (cache, TLB)
  - Prevent speculative code from causing higher costing misses (e.g. L2)
- Speculating through multiple branches
  - Complicates speculation recovery
  - No processor can resolve multiple branches per cycle




# Energy Efficiency

- Speculation and energy efficiency
  - Note: speculation is only energy efficient when it significantly improves performance
- Value prediction
  - Uses:
    - Loads that load from a constant pool
    - Instruction that produces a value from a small set of values
  - Not been incorporated into modern processors
  - Similar idea--*address aliasing prediction*--is used on some processors

# Branch Prediction

- Basic 2-bit predictor:
  - For each branch:
    - Predict taken or not taken
    - If the prediction is wrong two consecutive times, change prediction
- Correlating predictor:
  - Multiple 2-bit predictors for each branch
  - One for each possible combination of outcomes of preceding  $n$  branches
- Local predictor:
  - Multiple 2-bit predictors for each branch
  - One for each possible combination of outcomes for the last  $n$  occurrences of this branch
- Tournament predictor:
  - Combine correlating predictor with local predictor

# Branch Prediction Performance



Branch predictor performance